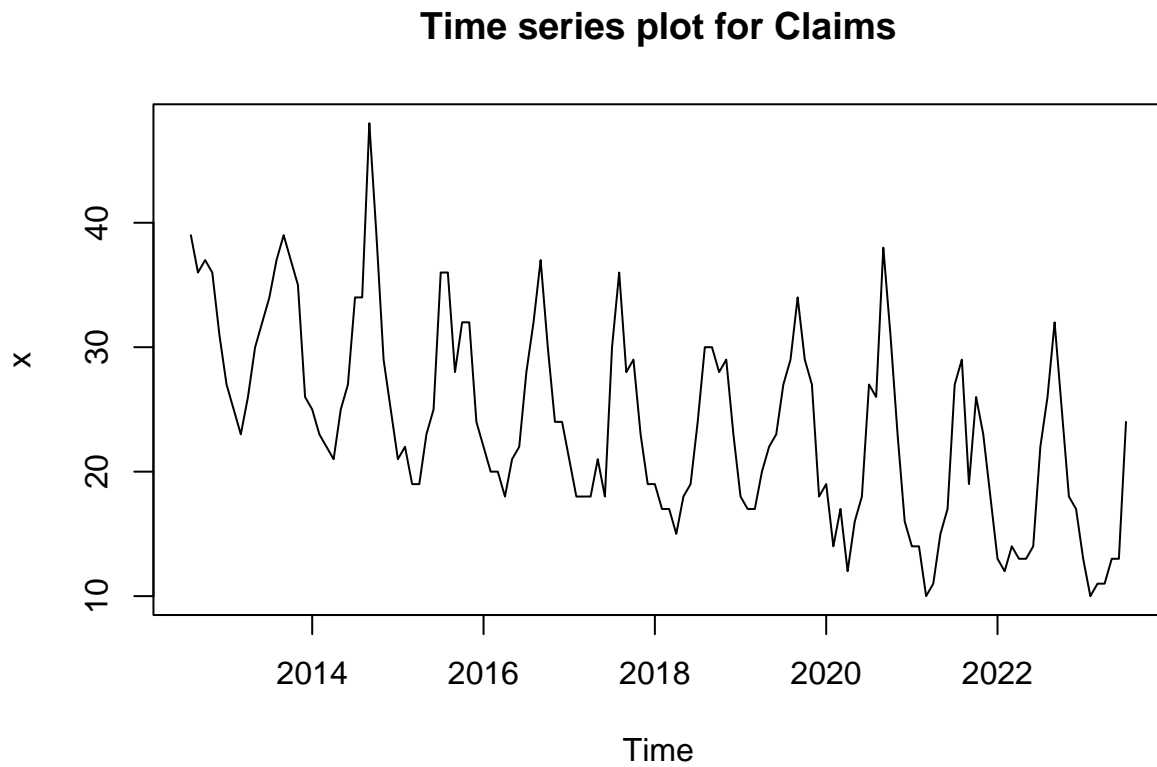


a)

```
plot(Claims, main = "Time series plot for Claims")
```




Comment:

- In the time series plot there seems to be seasonality. The period seems to be constant as a function of time which suggests the covariance among all claim observations is constant.
- The mean of the time series seems to be decreasing slightly with time. One could draw a line through the midpoints of the Y_t midpoints for some intervals of t to see this slightly negative slope.
- The variance looks like it could be constant with respect to time. However there are three peaks that stand out, suggesting the variance for these observations is unusually large.

b)

```
training_set <- window(Claims, start=start(Claims), end=time(Claims)[120])
test_set <- window(Claims, start=time(Claims)[121], end=end(Claims))
plot(training_set, col="blue", ylab = "Claims (Yt)", main = "Test and Training sets for time series of Claims")
lines(test_set, col="red")
legend(2012.4, 18, legend = c("Training set", "Test set"), col=c("blue", "red"), lty=1:1, cex=.5)
```



c)

```
tim <- time(training_set)
month <- as.factor(cycle(training_set))

reg_p1 <- lm(training_set ~ poly(as.vector(tim), 1) + month) # Orthogonal polynomial deg 1
reg_p2 <- lm(training_set ~ poly(as.vector(tim), 2) + month) # Orthogonal polynomial deg 2
reg_p3 <- lm(training_set ~ poly(as.vector(tim), 3) + month) # Orthogonal polynomial deg 3
reg_p4 <- lm(training_set ~ poly(as.vector(tim), 4) + month) # Orthogonal polynomial deg 4
reg_p5 <- lm(training_set ~ poly(as.vector(tim), 5) + month) # Orthogonal polynomial deg 5
reg_p6 <- lm(training_set ~ poly(as.vector(tim), 6) + month) # Orthogonal polynomial deg 6
reg_p7 <- lm(training_set ~ poly(as.vector(tim), 7) + month) # Orthogonal polynomial deg 7
reg_p8 <- lm(training_set ~ poly(as.vector(tim), 8) + month) # Orthogonal polynomial deg 8

tim.new <- as.vector(time(test_set))
month.new <- as.factor(cycle(test_set))
new <- data.frame(tim=tim.new, month=month.new)

pred_p1 <- predict.lm(reg_p1, new)
pred_p2 <- predict.lm(reg_p2, new)
pred_p3 <- predict.lm(reg_p3, new)
pred_p4 <- predict.lm(reg_p4, new)
pred_p5 <- predict.lm(reg_p5, new)
pred_p6 <- predict.lm(reg_p6, new)
pred_p7 <- predict.lm(reg_p7, new)
pred_p8 <- predict.lm(reg_p8, new)

# takes in a vector of ordered predictions
Y_test = as.vector(test_set)

mse_pred = function(Y_pred){
  mean((Y_test - Y_pred)^2)
}

msi_pred = function(Y_pred){
  mean((1-Y_test/Y_pred)^2)
}

p = seq(1,8)
MSEs = c(mse_pred(pred_p1),
          mse_pred(pred_p2),
          mse_pred(pred_p3),
          mse_pred(pred_p4),
          mse_pred(pred_p5),
          mse_pred(pred_p6),
          mse_pred(pred_p7),
          mse_pred(pred_p8))

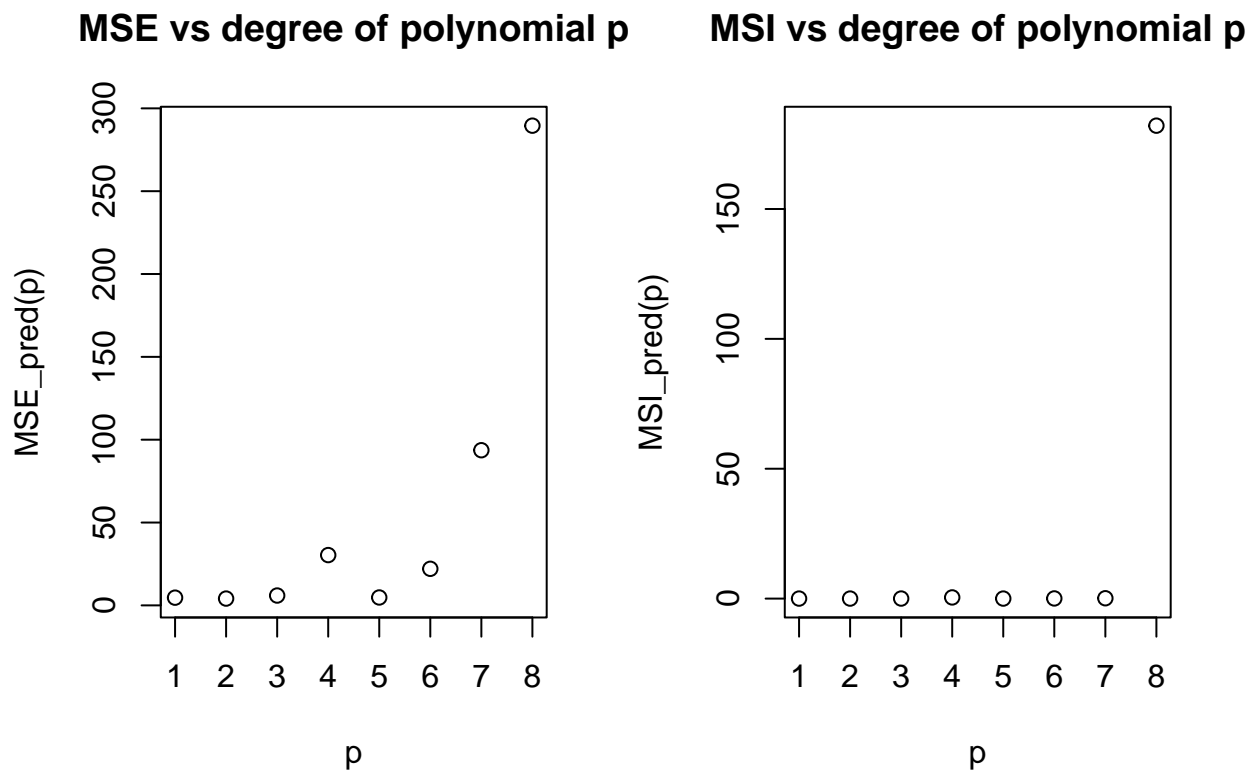
MSIs = c(msi_pred(pred_p1),
          msi_pred(pred_p2),
          msi_pred(pred_p3),
```

```

msi_pred(pred_p4),
msi_pred(pred_p5),
msi_pred(pred_p6),
msi_pred(pred_p7),
msi_pred(pred_p8))

par(mfrow=c(1,2))
plot(p, MSEs, ylab = "MSE_pred(p)",
     main = "MSE vs degree of polynomial p")
plot(p, MSIs, ylab = "MSI_pred(p)",
     main = "MSI vs degree of polynomial p")

```



Minimum value for MSEs is obtained on the second index, so the MSE divergence criteria picks $p=2$

```
min(MSEs)
```

```
## [1] 4.110931
```

```
MSEs
```

```
## [1] 4.646813 4.110931 5.946326 30.354078 4.772936 22.063624 93.645240
```

```
## [8] 289.542257
```

Minimum value for MSIs is obtained on the second index, so the MSI divergence criteria picks $p=2$

```
min(MSIs)
```

```
## [1] 0.009876646
```

```
MSIs
```

```
## [1] 1.093456e-02 9.876646e-03 1.790474e-02 4.458357e-01 1.104719e-02
```

```
## [6] 5.454036e-02 1.333643e-01 1.820798e+02
```


Both loss functions picked the same model with polynomial of degree two.

```
# using reg_p2 (regression polynomial degree 2)
reg_p2

##
## Call:
## lm(formula = training_set ~ poly(as.vector(tim), 2) + month)
##
## Coefficients:
##      (Intercept)  poly(as.vector(tim), 2)1  poly(as.vector(tim), 2)2
##           19.8470           -42.9129             3.5604
##           month2             month3             month4
##          -1.5869           -1.9744          -2.2625
##           month5             month6             month7
##           0.9487           2.1594           9.6694
##           month8             month9           month10
##          12.3255          13.3416          11.5571
##           month11           month12
##           7.9720           2.3863

# fitting the reg_p2 model to the Claims time
tim <- time(Claims)
t <- as.vector(tim)
t_ <- data.frame(tim=t, month=as.factor(cycle(Claims)))
mymodel <- predict.lm(reg_p2, t_)

plot(Claims)
# plotting predictions for reg_p2
lines(mymodel ~ t, type='l', col='red')
```



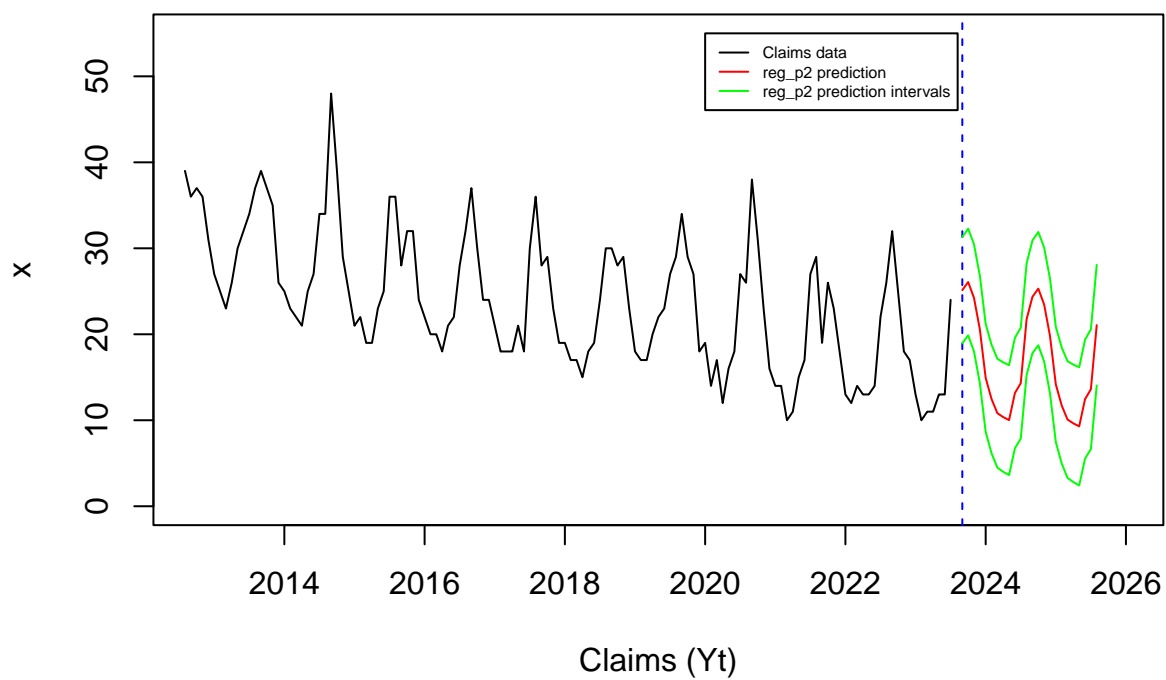
d)

```
tim.new <- seq(2023+8/12, 2025+7/12, by=1/12)
mons<- c(8:12,1:12, 1:7)
month.new <- factor(mons)
new <- data.frame(tim=tim.new,month=month.new)
# Forecast for the next 2 years using the chosen model reg_p2:
a <- predict.lm(reg_p2,new,interval='prediction', level=0.95)
# Numerical values for predictions along with prediction intervals:
a
```

##		fit	lwr	upr
## 1		25.133154	18.951175	31.31513
## 2		26.081007	19.875355	32.28666
## 3		24.228861	17.998752	30.45897
## 4		20.576714	14.321376	26.83205
## 5		14.924567	8.643236	21.20590
## 6		12.472420	6.164342	18.78050
## 7		10.820274	4.484702	17.15584
## 8		10.368127	4.004328	16.73193
## 9		10.015980	3.623228	16.40873
## 10		13.163834	6.741412	19.58625
## 11		14.311687	7.858891	20.76448
## 12		21.759540	15.275673	28.24341
## 13		24.353957	17.802957	30.90496
## 14		25.309079	18.720791	31.89737
## 15		23.464201	16.837745	30.09066
## 16		19.819323	13.153837	26.48481
## 17		14.174446	7.469080	20.87981
## 18		11.729568	4.983490	18.47565
## 19		10.084690	3.297082	16.87230
## 20		9.639812	2.809870	16.46975
## 21		9.294934	2.421870	16.16800
## 22		12.450056	5.533097	19.36702
## 23		13.605178	6.643564	20.56679
## 24		21.060300	14.053286	28.06731

```
plot(Claims,xlim=c(2012+8/12,2026), ylim=c(0, 55),
     xlab = "Claims (Yt)", main="Claims over time and prediction")
abline(v=2023+8/12,col='blue',lty=2) # adding a vertical line at the point where prediction starts
lines(a[,1]~tim.new,type='l',col='red')# plotting the predict
lines(a[,2]~tim.new,col='green') # plotting lower limit of the prediction interval
lines(a[,3]~tim.new,col='green') # plotting upper limit of the prediction interval
legend(2020, 55, legend = c("Claims data","reg_p2 prediction", "reg_p2 prediction intervals"), col=c("b",
```

Claims over time and prediction



e) Residual Diagnostics

We must test residuals for normality, zero mean, constant variance, and randomness (iid). Note since we used the training set to train the model, we will use only the residuals of the training set.

Normality test

```
# Model used
mymodel <- reg_p2
# Testing normality with Shapiro-Wilk
residuals.reg = residuals(mymodel) #extracting the residuals
shapiro.test(residuals.reg) #Shapiro-Wilk test of normality H0: residuals are normal
```


```
##
##  Shapiro-Wilk normality test
##
## data:  residuals.reg
## W = 0.93537, p-value = 2.114e-05
```

The Null hypothesis for the Shapiro-Wilk test is that the residuals are normally distributed. the small p -value suggests there is strong evidence against this. So this assumption is **not satisfied**

Zero-mean

```
plot(residuals.reg, ylab = "Residuals for reg_2 model", main = "Residuals plot")
abline(h=0, col="red", lty=2, lwd=2)
legend(0,10,legend=c("zero-line"), col="red", lty=2, cex=.75)
```


Residuals plot



From visual inspection, the residuals seem to have zero mean as they look evenly spread around the zero line. So this assumption is **satisfied**.

Constant variance:

```
# Testing Constant variance with FLigner-Killeen
#First, 6 groups of 20 obs.
plot(residuals.reg, pch=16, col=adjustcolor("black",0.5)) # plotting the residuals vs time
abline(v=c(1:5)*20, col="red", lwd=2, lty=2)
```



```
segments = factor(rep(1:6,each=20)) #creating 6 chunks to test variance homogeneity
fligner.test(residuals.reg, segments) # Fligner's test for H0:sigma1=sigma2=...=sigma6 (note that we m
```


```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: residuals.reg and segments
## Fligner-Killeen:med chi-squared = 6.7425, df = 5, p-value = 0.2405
```

The p -value is high at 0.2405, so there is no evidence against the null hypothesis of equal variance. we conclude the constant variance assumption is **satisfied**.

Randomness (iid)

```
# Testing randomness with difference sign test and runs test
randtests::difference.sign.test(residuals.reg)
```

```
##
## Difference Sign Test
##
## data: residuals.reg
## statistic = -0.15746, n = 120, p-value = 0.8749
## alternative hypothesis: nonrandomness
randtests::runs.test(residuals.reg, plot = TRUE)
```



```
##
## Runs Test
##
## data: residuals.reg
## statistic = -1.4668, runs = 53, n1 = 60, n2 = 60, n = 120, p-value =
## 0.1424
## alternative hypothesis: nonrandomness
```

The *Difference Sign Test* has a large p -value at 0.8749 which means there was no evidence to reject the null hypothesis of random residuals. Therefore under this test, the randomness assumption **is satisfied**.

The *Runs test* has a p -value smaller than that of the DST but still large enough to suggest there is no evidence against the assumption of randomness. Therefore under this test, the randomness assumption **is satisfied**.

The normality assumption was not satisfied, so we should not make inferences with this model. Maybe trying make the mean stationary would solve this.